

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df=pd.read_csv('/content/Fuel_Consumption_2000-2022[1].csv')
df.head(20)
```

| | YEAR | MAKE | MODEL | VEHICLE CLASS | ENGINE SIZE | CYLINDERS | TRANSMISSION | FUEL | FUEL CONSUMPTION | HWY (L/100 km) |
|----|------|-------|-------------------|---------------|-------------|-----------|--------------|------|------------------|----------------|
| 0 | 2000 | ACURA | 1.6EL | COMPACT | 1.6 | 4 | A4 | X | 9.2 | 6.7 |
| 1 | 2000 | ACURA | 1.6EL | COMPACT | 1.6 | 4 | M5 | X | 8.5 | 6.5 |
| 2 | 2000 | ACURA | 3.2TL | MID-SIZE | 3.2 | 6 | AS5 | Z | 12.2 | 7.4 |
| 3 | 2000 | ACURA | 3.5RL | MID-SIZE | 3.5 | 6 | A4 | Z | 13.4 | 9.2 |
| 4 | 2000 | ACURA | INTEGRA | SUBCOMPACT | 1.8 | 4 | A4 | X | 10.0 | 7.0 |
| 5 | 2000 | ACURA | INTEGRA | SUBCOMPACT | 1.8 | 4 | M5 | X | 9.3 | 6.8 |
| 6 | 2000 | ACURA | INTEGRA GSR/TYPED | SUBCOMPACT | 1.8 | 4 | M5 | Z | 9.4 | 7.0 |
| 7 | 2000 | ACURA | NSX | SUBCOMPACT | 3.0 | 6 | AS4 | Z | 13.6 | 9.2 |
| 8 | 2000 | ACURA | NSX | SUBCOMPACT | 3.2 | 6 | M6 | Z | 13.8 | 9.1 |
| 9 | 2000 | AUDI | A4 | COMPACT | 1.8 | 4 | A5 | Z | 11.4 | 7.2 |
| 10 | 2000 | AUDI | A4 | COMPACT | 1.8 | 4 | M5 | Z | 9.7 | 6.8 |
| 11 | 2000 | AUDI | A4 | COMPACT | 2.8 | 6 | A5 | Z | 13.0 | 8.2 |
| 12 | 2000 | AUDI | A4 | COMPACT | 2.8 | 6 | M5 | Z | 11.7 | 7.5 |
| 13 | 2000 | AUDI | A4 QUATTRO | COMPACT | 1.8 | 4 | A5 | Z | 12.1 | 7.7 |
| 14 | 2000 | AUDI | A4 QUATTRO | COMPACT | 1.8 | 4 | M5 | Z | 10.7 | 7.5 |
| 15 | 2000 | AUDI | A4 QUATTRO | COMPACT | 2.8 | 6 | A5 | Z | 13.3 | 8.5 |
| 16 | 2000 | AUDI | A4 QUATTRO | COMPACT | 2.8 | 6 | M5 | Z | 12.7 | 8.7 |

```
df.tail(15)
```

Data processing

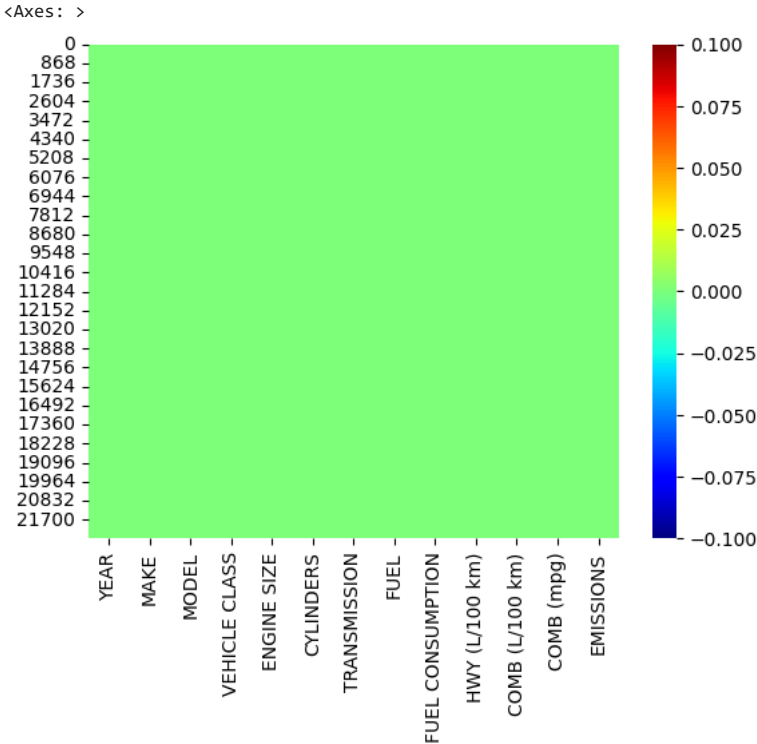
```
df.shape

(22556, 13)
2014 2015 2016 2017 2018 2019 2020 2021 2022 2023 2024 2025 2026
volkswagen 4motion Small 2.0 4 A30 A 10.0 0.0

#check for null values
df.isna().sum()

YEAR 0
MAKE 0
MODEL 0
VEHICLE CLASS 0
ENGINE SIZE 0
CYLINDERS 0
TRANSMISSION 0
FUEL 0
FUEL CONSUMPTION 0
HWY (L/100 km) 0
COMB (L/100 km) 0
COMB (mpg) 0
EMISSIONS 0
dtype: int64


#plot heatmap for null values
sns.heatmap(df.isnull(), cmap='jet')
```



```
#correlation
corr=df.corr()
corr
```

```
<invthon-inout-169-78aeh47f0e90>: FutureWarning: The default value of numeric_only in DataFrame.corr
#heatmap of correlation
sns.heatmap(corr, cmap='viridis')
```

<Axes: >



```
df.dtypes
```

```
YEAR                int64
MAKE                object
MODEL               object
VEHICLE CLASS       object
ENGINE SIZE         float64
CYLINDERS           int64
TRANSMISSION        object
FUEL                object
FUEL CONSUMPTION     float64
HWY (L/100 km)      float64
COMB (L/100 km)     float64
COMB (mpg)          int64
EMISSIONS           int64
dtype: object
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22556 entries, 0 to 22555
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   YEAR                  22556 non-null int64
1   MAKE                  22556 non-null object
2   MODEL                 22556 non-null object
3   VEHICLE CLASS        22556 non-null object
4   ENGINE SIZE           22556 non-null float64
5   CYLINDERS             22556 non-null int64
6   TRANSMISSION          22556 non-null object
7   FUEL                  22556 non-null object
8   FUEL CONSUMPTION      22556 non-null float64
9   HWY (L/100 km)       22556 non-null float64
10  COMB (L/100 km)       22556 non-null float64
11  COMB (mpg)            22556 non-null int64
12  EMISSIONS             22556 non-null int64
dtypes: float64(4), int64(4), object(5)
memory usage: 2.2+ MB
```

```
#I just renamed the column 'MAKE' to 'Company' for better understanding
df.rename(columns={'MAKE': 'Company'}, inplace=True)
df
```

| | YEAR | Company | MODEL | VEHICLE CLASS | ENGINE SIZE | CYLINDERS | TRANSMISSION | FUEL | FUEL CONSUMPTION | HWY (L/100 km) |
|-------|------|---------|-------------|---------------|-------------|-----------|--------------|------|------------------|----------------|
| 0 | 2000 | ACURA | 1.6EL | COMPACT | 1.6 | 4 | A4 | X | 9.2 | 6 |
| 1 | 2000 | ACURA | 1.6EL | COMPACT | 1.6 | 4 | M5 | X | 8.5 | 6 |
| 2 | 2000 | ACURA | 3.2TL | MID-SIZE | 3.2 | 6 | AS5 | Z | 12.2 | 7 |
| 3 | 2000 | ACURA | 3.5RL | MID-SIZE | 3.5 | 6 | A4 | Z | 13.4 | 9 |
| 4 | 2000 | ACURA | INTEGRA | SUBCOMPACT | 1.8 | 4 | A4 | X | 10.0 | 7 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 22551 | 2022 | Volvo | XC40 T5 AWD | SUV: Small | 2.0 | 4 | AS8 | Z | 10.7 | 7 |
| 22552 | 2022 | Volvo | XC60 B5 AWD | SUV: Small | 2.0 | 4 | AS8 | Z | 10.5 | 8 |

```
#unique values in each column]
cols=df.columns
for cols in (df.columns):
    print ("Unique columns of", cols, "\n", df[cols].unique())
    print ("-----")

    28.1 27.3 28.3 26.5]
    -----
Unique columns of HWY (L/100 km)
[ 6.7  6.5  7.4  9.2  7.   6.8  9.1  7.2  8.2  7.5  7.7  8.5  8.7  8.6
  9.   7.6  7.3  8.9  7.8  7.9  8.4  8.3  8.8 10.1 10.8 10.4  8.  11.1
  8.1 13.4 11.4 11.9  9.9 12.4 12.3 11.  12.  13.6  7.1  6.9 11.8 12.7
 12.2 12.5  4.5  6.2  4.9 12.6  9.7 10.9  9.5  5.9  6.4  6.  12.1 10.2
  9.6 13.2 10.7 10.3 11.7 11.5 13.3 13.  12.9 11.6 13.7 14.  13.5 14.5
 16.6 11.3  6.6 13.1 11.2 10.6  9.3  9.4 10.5  6.1  5.8  3.2 10.   9.8
 14.7  5.4  5.7  5.3  6.3  5.6  5.5  4.8  4.4 12.8 13.8 14.9 15.4  5.2
  4.7 16.4 15.9 15.7 16.9 15.8 16.2 16.3 14.4 14.3 14.8 16.1 17.7 16.
 15.  17.6 15.5 15.1  4.6 13.9  5.1  4.3  5.  14.1 14.6 14.2  3.3  4.2
 17.9 19.   3.8 15.6 16.5 17.3 17.4 17.  15.2 15.3 16.8 17.1  4.  16.7
 17.5 17.2 20.5 20.6 18.6 18.5 18.  18.1 20.9  4.1  3.9 19.6]
    -----
Unique columns of COMB (L/100 km)
[ 8.1  7.6 10.   11.5  8.6  8.2  8.3 11.6 11.7  9.5  8.4 10.8  9.8 10.1
  9.3 11.1 10.9 11.2 11.3 10.7 10.3  9.7 10.4 12.2 14.4 10.6 14.6 10.5
 11.   9.9 16.6 13.4 13.8 12.8 14.1 12.5 14.9 14.7 12.7 13.1 14.3 16.8
 10.2  8.7  8.9  9.4  8.8  9.1 14.2 15.3 16.4 14.  15.6  5.3  7.1  5.7
 15.  15.1 12.   12.3 13.2 12.6  9.6  8.5  7.4  7.7  9.  15.7 13.6 13.9
 16.3 14.8 17.2 11.8 12.4 13.3 15.2 16.1 16.9 14.5 18.4 18.3 19.9 19.2
 22.7 11.4 15.5 12.9 13.7 13.5 13.  11.9 16.7  7.5  6.8  3.6  7.8  7.3
  9.2 12.1 15.4 17.7  7.9  6.7  7.   6.6  6.3  6.   5.1 15.8 16.  16.5
 17.  16.2 17.8  8.  18.5 17.5 17.4  7.2  6.4  6.1  4.6  6.9 20.7 17.9
 18.2 20.9 18.6 19.  19.1 17.1 17.3 20.6  6.5 19.3 19.7 20.3 19.8 24.8
  4.8  5.9  5.  19.5 18.   4.1  5.5  5.6 17.6 18.7 21.3 23.2 22.3 20.4
 18.1 19.4  4.2 20.  15.9  4.5 18.9 19.6 20.1 21.4 21.7  5.8 18.8 20.5
  5.4 22.1 20.2  4.7 21.   3.8  6.2 23.  23.1 20.8 21.5 23.3  4.3  4.9
  3.7  5.2  4.4 25.9 26.1 22.2 22.9  4.  22.4 21.2]
    -----
Unique columns of COMB (mpg)
[35 37 28 25 33 34 24 30 26 29 27 23 20 19 17 21 22 32 31 18 53 40 50 38
 16 15 14 12 42 78 36 39 43 45 47 55 44 46 61 41 11 59 48 56 69 51 13 67
 63 49 52 60 74 66 58 76 54 64 71]
    -----
Unique columns of EMISSIONS
[186 175 230 264 198 189 191 267 269 218 193 248 225 232 214 255 251 258
 260 246 237 223 239 281 331 244 336 242 253 228 382 308 317 294 324 288
 343 338 292 301 329 386 235 200 205 216 202 209 327 352 377 322 359 122
 163 131 345 347 276 283 304 290 221 196 170 177 207 313 320 375 340 396
 271 285 306 350 370 389 334 314 423 421 458 442 522 262 356 297 315 310
 299 274 384 361 172 156  83 179 168 212 278 354 407 182 154 161 152 145
 162 138 363 243 368 380 391 316 319 373 184 426 402 400 166 147 140 106
 159 409 318 412 344 428 393 398 321 330 150 444 467 208 455 570 312 259
 227 110 136 165 135 307 291 346  94 148 151 405 270 341 296 371 513 469
 419 453 414 275 229 265 113 272 197 286 366 104 430 173 305 337 222 233
 339 238 282 302 326 277 298 266 342 416 289 181 446 490 335 178 133 219
 203 293 328 273 124 174 280 254 210 508 323 187 261 245 115 309 108  87
 143 240 325 333 126 195 439 256  99 211 117  85 171 224 139 146 141 120
 194 226 101 234 435 357 142 311 127 379 217 388 213 192 432 437 418 250
 287 206 129 183 180 176 185 249 365 332 362 390 204 201 247 257 263 300
 360 408 231 241 284 215 199 268 252 220 279 167 160 190 236 441 465 295
 188 157 353 364 303 153 130 155 169 158 420 452 132 417 476 369 355 401
 348 403 404 367 134 137 111 351 450 349 438 445 121 387 164 406 114 103
  96 461 372 358 413 128 149 105 378 118 102 454 464 473 410 397 383 381
 487 493 109 537 392 485 535 395 385 608 520 515 539 374 489 498]
    -----

#value counts of each column
for cols in (df.columns):
```

```
print ("Value counts of", cols, "\n", df[cols].value_counts())
print ("-----")
```

```
35    460
17    460
36    414
37    344
38    335
16    291
39    264
15    251
40    246
42    206
14    170
43    142
41    114
13     64
46     61
45     59
47     52
44     47
50     44
55     29
48     29
12     27
53     23
63     20
49     19
58     18
51     17
56     15
60     14
59     14
61     14
52     12
54     11
69     10
11      8
78      7
74      7
67      6
64      6
66      4
76      2
71      2
```

```
Name: COMB (mpg), dtype: int64
```

```
-----
Value counts of EMISSIONS
```

```
221    364
225    352
228    348
232    339
230    338
...
372     1
134     1
351     1
450     1
417     1
```

```
Name: EMISSIONS, Length: 358, dtype: int64
```

```
-----
```

```
#from below feature selection step, model and transmission columns are having low values/prior
df1=df.drop(['MODEL','TRANSMISSION'], axis=1)
df1
```

```
#I applied get_dummies to columns 'Company', and 'Vehicle class for better visibility
dummy=pd.get_dummies(df1[['Company','VEHICLE CLASS']], drop_first=True)
dummy
```

| | Company_ALFA ROMEO | Company_ASTON MARTIN | Company_AUDI | Company_Acura | Company_Alfa Romeo | Company_Aston Martin | Company_A |
|-------|-----------------------|-------------------------|--------------|---------------|-----------------------|-------------------------|-----------|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 2 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 4 | 0 | 0 | 0 | 0 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | |
| 22551 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 22552 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 22553 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 22554 | 0 | 0 | 0 | 0 | 0 | 0 | |
| 22555 | 0 | 0 | 0 | 0 | 0 | 0 | |

22556 rows × 117 columns



```
#combine both dummy and main df
dfe=pd.concat([df1,dummy],axis=1)
dfe
```

| | YEAR | Company | VEHICLE CLASS | ENGINE SIZE | CYLINDERS | FUEL | FUEL CONSUMPTION | HWY (L/100 km) | COMB (L/100 km) | COMB (mpg) | ... | C |
|-------|------|---------|------------------|----------------|-----------|------|---------------------|----------------------|-----------------------|---------------|-----|---|
| 0 | 2000 | ACURA | COMPACT | 1.6 | 4 | X | 9.2 | 6.7 | 8.1 | 35 | ... | |
| 1 | 2000 | ACURA | COMPACT | 1.6 | 4 | X | 8.5 | 6.5 | 7.6 | 37 | ... | |
| 2 | 2000 | ACURA | MID-SIZE | 3.2 | 6 | Z | 12.2 | 7.4 | 10.0 | 28 | ... | |
| 3 | 2000 | ACURA | MID-SIZE | 3.5 | 6 | Z | 13.4 | 9.2 | 11.5 | 25 | ... | |
| 4 | 2000 | ACURA | SUBCOMPACT | 1.8 | 4 | X | 10.0 | 7.0 | 8.6 | 33 | ... | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 22551 | 2022 | Volvo | SUV: Small | 2.0 | 4 | Z | 10.7 | 7.7 | 9.4 | 30 | ... | |
| 22552 | 2022 | Volvo | SUV: Small | 2.0 | 4 | Z | 10.5 | 8.1 | 9.4 | 30 | ... | |
| 22553 | 2022 | Volvo | SUV: Small | 2.0 | 4 | Z | 11.0 | 8.7 | 9.9 | 29 | ... | |
| 22554 | 2022 | Volvo | SUV: Standard | 2.0 | 4 | Z | 11.5 | 8.4 | 10.1 | 28 | ... | |
| 22555 | 2022 | Volvo | SUV: Standard | 2.0 | 4 | Z | 12.4 | 8.9 | 10.8 | 26 | ... | |

22556 rows × 128 columns



```
#after combining, drop the unwanted columns
dfe=dfe.drop(['Company','VEHICLE CLASS'], axis=1)
dfe
```

| | YEAR | ENGINE SIZE | CYLINDERS | FUEL | FUEL CONSUMPTION | HWY (L/100 km) | COMB (L/100 km) | COMB (mpg) | EMISSIONS | Company_ALFA ROMEO | ... |
|-------|------|-------------|-----------|------|------------------|----------------|-----------------|------------|-----------|--------------------|-----|
| 0 | 2000 | 1.6 | 4 | X | 9.2 | 6.7 | 8.1 | 35 | 186 | 0 | ... |
| 1 | 2000 | 1.6 | 4 | X | 8.5 | 6.5 | 7.6 | 37 | 175 | 0 | ... |
| 2 | 2000 | 3.2 | 6 | Z | 12.2 | 7.4 | 10.0 | 28 | 230 | 0 | ... |
| 3 | 2000 | 3.5 | 6 | Z | 13.4 | 9.2 | 11.5 | 25 | 264 | 0 | ... |
| 4 | 2000 | 1.8 | 4 | X | 10.0 | 7.0 | 8.6 | 33 | 198 | 0 | ... |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 22551 | 2022 | 2.0 | 4 | Z | 10.7 | 7.7 | 9.4 | 30 | 219 | 0 | ... |
| 22552 | 2022 | 2.0 | 4 | Z | 10.5 | 8.1 | 9.4 | 30 | 219 | 0 | ... |

dfe.columns

```
Index(['YEAR', 'ENGINE SIZE', 'CYLINDERS', 'FUEL', 'FUEL CONSUMPTION',
      'HWY (L/100 km)', 'COMB (L/100 km)', 'COMB (mpg)', 'EMISSIONS',
      'Company_ALFA ROMEO',
      ...
      'VEHICLE CLASS_SUV: Standard', 'VEHICLE CLASS_Special purpose vehicle',
      'VEHICLE CLASS_Station wagon: Mid-size',
      'VEHICLE CLASS_Station wagon: Small', 'VEHICLE CLASS_Subcompact',
      'VEHICLE CLASS_TWO-SEATER', 'VEHICLE CLASS_Two-seater',
      'VEHICLE CLASS_VAN - CARGO', 'VEHICLE CLASS_VAN - PASSENGER',
      'VEHICLE CLASS_Van: Passenger'],
      dtype='object', length=126)

#label encode the data's to convert it to numeric
nw_cols=dfe.columns
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
for nw_cols in dfe.columns:
    dfe[nw_cols]=le.fit_transform(dfe[nw_cols])
```

dfe.dtypes

```
YEAR          int64
ENGINE SIZE    int64
CYLINDERS      int64
FUEL           int64
FUEL CONSUMPTION int64
...
VEHICLE CLASS_TWO-SEATER int64
VEHICLE CLASS_Two-seater int64
VEHICLE CLASS_VAN - CARGO int64
VEHICLE CLASS_VAN - PASSENGER int64
VEHICLE CLASS_Van: Passenger int64
Length: 126, dtype: object
```

dfe

```

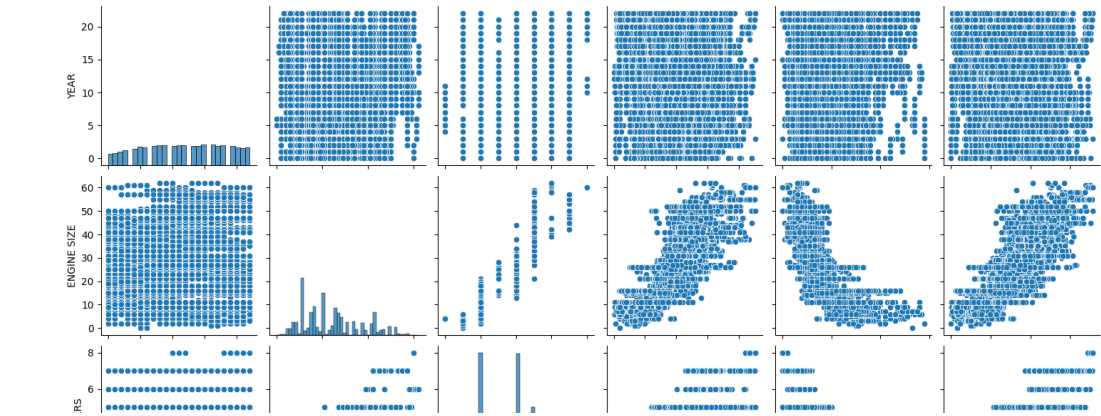
    HWY      COMB      COMB      EMISSIONS      Company_ALFA      VEHICLE
    / 1100  / 1100      COMB      EMISSIONS      Company_ALFA      VEHICLE
dfe['FUEL CONSUMPTION'].value_counts()

85      399
82      360
71      346
80      341
81      339
...
218      1
212      1
15      1
1      1
214      1
Name: FUEL CONSUMPTION, Length: 228, dtype: int64
,      45      57      10      117      0      0      0      0
```

Visualization

```


#pairplot of columns=('YEAR', 'ENGINE SIZE', 'CYLINDERS', 'FUEL CONSUMPTION', 'COMB (mpg)', 'EMISSIONS')
sns.pairplot(dfe[['YEAR', 'ENGINE SIZE', 'CYLINDERS', 'FUEL CONSUMPTION', 'COMB (mpg)', 'EMISSIONS']])
plt.show()
```


```
dfe['FUEL'].value_counts()
```

```
3    11822
4     9316
1     1071
0       314
2        33
Name: FUEL, dtype: int64
```


```
#Distribution of fuel type
plt.figure(figsize=(8, 5))
sns.countplot(x='FUEL', data=dfe)
plt.title("Distribution of Fuel Types")
plt.xlabel("Fuel Type")
plt.ylabel("Count")
plt.show()
#where,
#0-D
#1-E
#2-N
#3-X
#4-Z
```



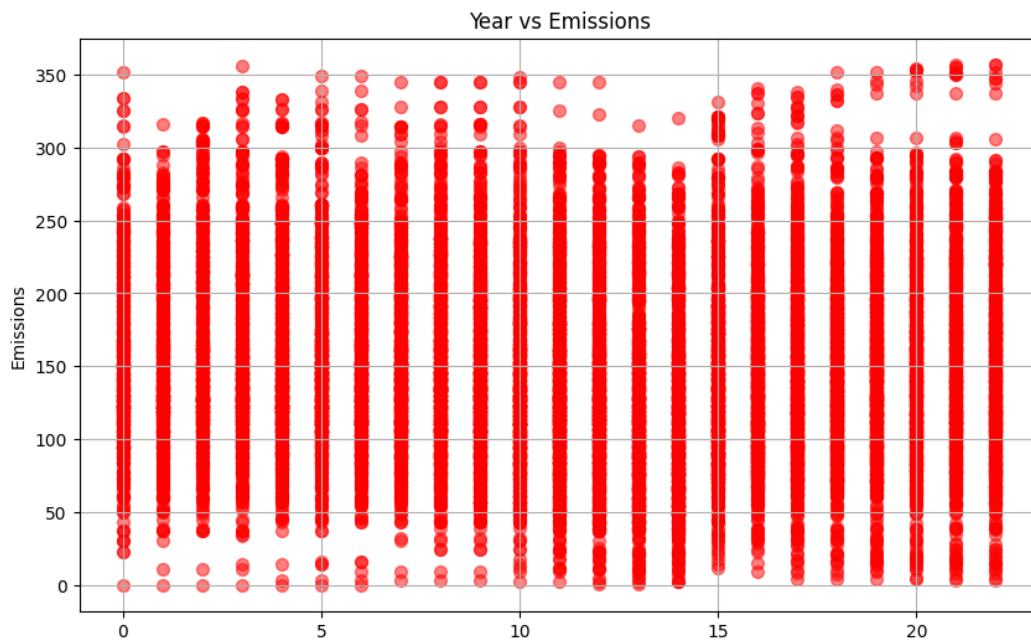
```
#fuel consumption by fuel type
plt.figure(figsize=(8, 5))
sns.boxplot(x='FUEL', y='FUEL CONSUMPTION', data=dfe)
plt.title("Fuel Consumption Distribution by Fuel Type")
plt.xlabel("Fuel Type")
plt.ylabel("Fuel Consumption (L/100 km)")
plt.show()
```



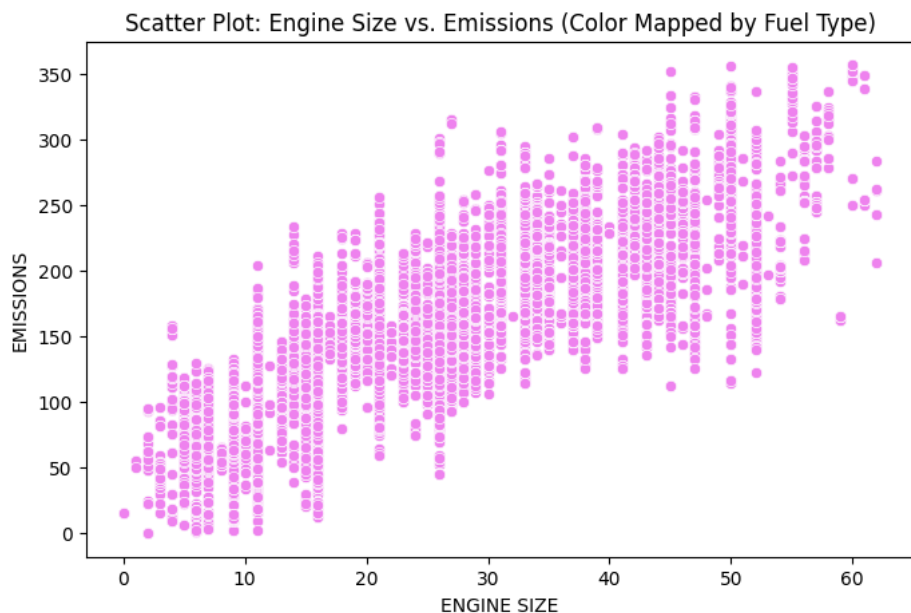
```
#distribution of fuel types
fuel_counts=dfe['FUEL'].value_counts()
plt.figure(figsize=(8, 8))
plt.pie(fuel_counts.values, labels=fuel_counts.index, autopct='%1.1f%%', startangle=140)
plt.title("Distribution of Fuel Types")
plt.axis('equal')
plt.show()
#where,
#0-D
#1-E
#2-N
#3-X
#4-Z
```



```
#emission by year
plt.figure(figsize=(10, 6))
plt.scatter(dfe['YEAR'], dfe['EMISSIONS'], s=50, alpha=0.5, color='red')
plt.xlabel('Year')
plt.ylabel('Emissions')
plt.title('Year vs Emissions')
plt.grid(True)
plt.show()
```




```
plt.figure(figsize=(8, 5))
sns.scatterplot(x='ENGINE SIZE', y='EMISSIONS', data=dfe, color='violet')
plt.title("Scatter Plot: Engine Size vs. Emissions (Color Mapped by Fuel Type)")
plt.show()
```




```
#cylinders by emissions
plt.figure(figsize=(8, 5))
sns.scatterplot(x='CYLINDERS', y='EMISSIONS', data=dfe, color='brown')
plt.title('Cylinders by emissions')
```

```
Text(0.5, 1.0, 'Cylinders by emissions')
```



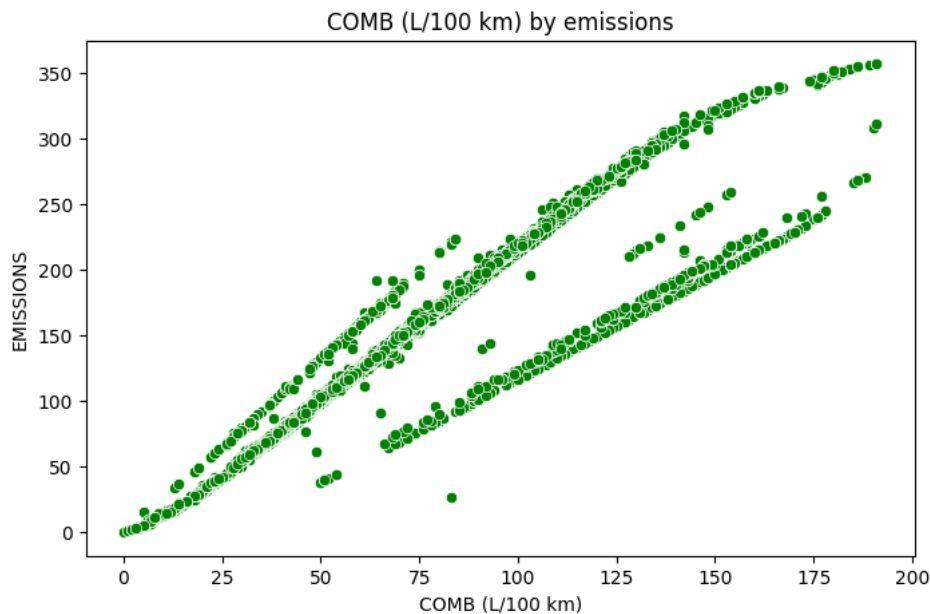
```
#HWY (L/100 km) by emissions
plt.figure(figsize=(8, 5))
sns.scatterplot(x='HWY (L/100 km)', y='EMISSIONS', data=dfe)
plt.title('HWY (L/100 km) by emissions')
```

```
Text(0.5, 1.0, 'HWY (L/100 km) by emissions')
```




```
#COMB (L/100 km) by emissions
plt.figure(figsize=(8, 5))
sns.scatterplot(x='COMB (L/100 km)', y='EMISSIONS', data=dfe, color='green')
plt.title('COMB (L/100 km) by emissions')
```

```
Text(0.5, 1.0, 'COMB (L/100 km) by emissions')
```





```
#Different cylinders
cylinder_counts=dfe['CYLINDERS'].value_counts()
plt.figure(figsize=(8, 6))
plt.bar(cylinder_counts.index, cylinder_counts.values)
plt.title("Distribution of Cylinders")
plt.xlabel("Number of Cylinders")
plt.ylabel("Count")
plt.show()
```





```
#Emission by each year
sns.catplot(data=dfe, x='FUEL', y='EMISSIONS', kind='bar', col='YEAR', col_wrap=2, sharex=False)
plt.xticks(rotation='vertical')
```

```
/usr/local/lib/python3.10/dist-packages/seaborn/categorical.py:3201: UserWarning: Setting `sharex=False`
warnings.warn(msg.format("sharex", "x"), UserWarning)
(array([0, 1, 2, 3]),
 [Text(0, 0, '0'), Text(1, 0, '1'), Text(2, 0, '3'), Text(3, 0, '4')])
```










```
#top 10 car company contribution in the market
plt.figure(figsize=(10,10))
Top_10_cars=df1['Company'].value_counts().head(10)
plt.pie(Top_10_cars, labels=Top_10_cars.index, autopct="%0.1f%%")
plt.title('Top Ten Car Company Contribution in the market')
```


```
Text(0.5, 1.0, 'Top Ten Car Company Contribution in the market')
```

Top Ten Car Company Contribution in the market



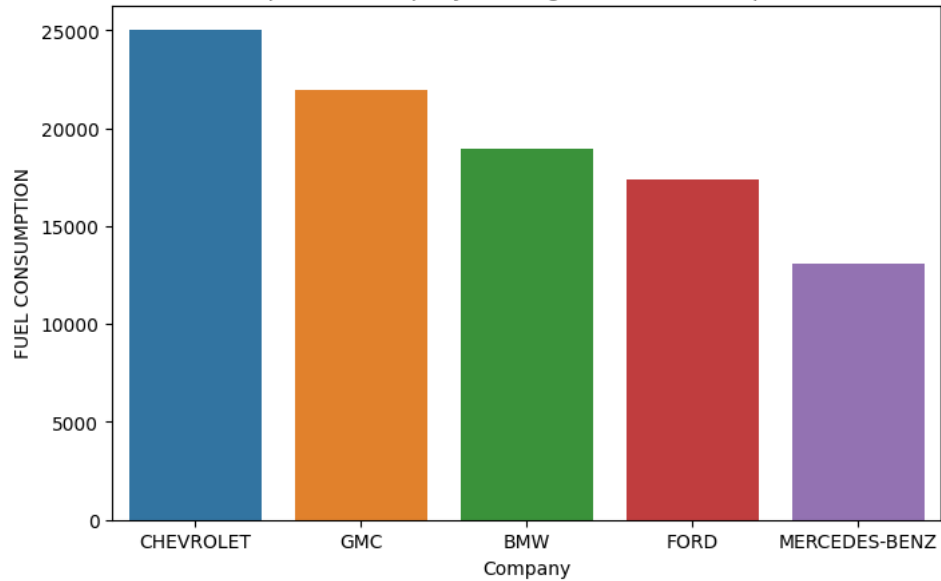
```
#top 10 vehicle class in the market
plt.figure(figsize=(10,10))
Top_10_vehicleclass=df1['VEHICLE CLASS'].value_counts().head(10)
plt.pie(Top_10_vehicleclass, labels=Top_10_vehicleclass.index, autopct="%0.1f%%")
plt.title('Top Ten Car Company Contribution in the market')
```

```
Text(0.5, 1.0, 'Top Ten Car Company Contribution in the market')
Top Ten Car Company Contribution in the market
```



```
#top 5 cars company with highest fuel consumption
company_consumption=df1.groupby('Company')['FUEL CONSUMPTION'].sum().sort_values(ascending = False).reset_index()
plt.figure(figsize=(8,5))
sns.barplot(x="Company", y="FUEL CONSUMPTION", data=company_consumption[:5])
plt.title('Top 5 cars company with highest fuel consumption')
```

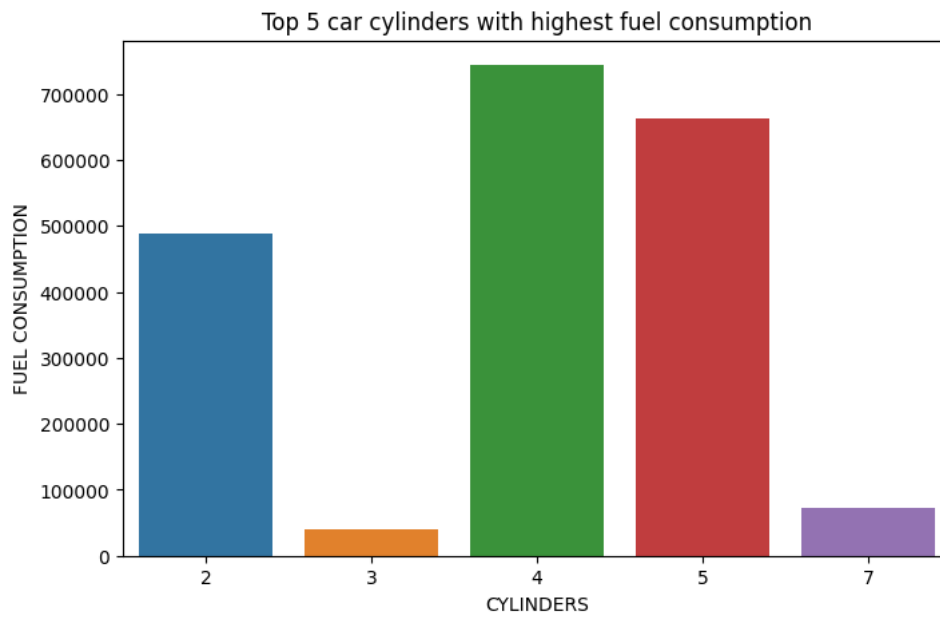
```
Text(0.5, 1.0, 'Top 5 cars company with highest fuel consumption')
Top 5 cars company with highest fuel consumption
```



```
#top 5 car cylinders with highest fuel consumption
carcyl_consumption=dfe.groupby('CYLINDERS')['FUEL CONSUMPTION'].sum().sort_values(ascending = False).reset_index()
plt.figure(figsize=(8,5))
```


```
sns.barplot(x="CYLINDERS", y="FUEL_CONSUMPTION", data=carcyl_consumption[:5])
plt.title('Top 5 car cylinders with highest fuel consumption')
```

```
Text(0.5, 1.0, 'Top 5 car cylinders with highest fuel consumption')
```



```
#top 5 years with high fuel consumption
plt.figure(figsize=(10, 6))
top_fuel=dfe.groupby('YEAR')['FUEL_CONSUMPTION'].sum().sort_values(ascending = False).index
plt.pie(top_fuel[:5],labels=top_fuel[:5], autopct="%1.2f%%")
plt.title('top 5 years with high fuel consumption')
```


```
Text(0.5, 1.0, 'top 5 years with high fuel consumption')
```



Double-click (or enter) to edit

```
#Top 5 fuel with high fuel consumption
top_5_fuels=dfe.groupby('FUEL')['FUEL_CONSUMPTION'].sum().sort_values(ascending=False).index
plt.figure(figsize=(10, 6))
plt.pie(top_5_fuels[:5], labels=top_5_fuels[:5], autopct="%1.2f%%")
plt.title("Top 5 full with high fuel consumption")
#where,
#0-D
#1-E
#2-N
#3-X
#4-Z
```

Text(0.5, 1.0, 'Top 5 full with high fuel consumption')
Top 5 full with high fuel consumption



```
x=dfe.drop(['FUEL CONSUMPTION'], axis=1)
```

x

| | YEAR | ENGINE SIZE | CYLINDERS | FUEL | HWY (L/100 km) | COMB (L/100 km) | COMB (mpg) | EMISSIONS | Company_ALFA ROMEO | Company_ASTON MARTIN | ... |
|-------|------|-------------|-----------|------|----------------|-----------------|------------|-----------|--------------------|----------------------|-----|
| 0 | 0 | 7 | 2 | 3 | 31 | 44 | 24 | 84 | 0 | 0 | .. |
| 1 | 0 | 7 | 2 | 3 | 29 | 39 | 26 | 73 | 0 | 0 | .. |
| 2 | 0 | 23 | 4 | 4 | 38 | 63 | 17 | 128 | 0 | 0 | .. |
| 3 | 0 | 26 | 4 | 4 | 56 | 78 | 14 | 162 | 0 | 0 | .. |
| 4 | 0 | 9 | 2 | 3 | 34 | 49 | 22 | 96 | 0 | 0 | .. |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | .. |
| 22551 | 22 | 11 | 2 | 4 | 41 | 57 | 19 | 117 | 0 | 0 | .. |
| 22552 | 22 | 11 | 2 | 4 | 45 | 57 | 19 | 117 | 0 | 0 | .. |
| 22553 | 22 | 11 | 2 | 4 | 51 | 62 | 18 | 130 | 0 | 0 | .. |
| 22554 | 22 | 11 | 2 | 4 | 48 | 64 | 17 | 134 | 0 | 0 | .. |
| 22555 | 22 | 11 | 2 | 4 | 53 | 71 | 15 | 150 | 0 | 0 | .. |

22556 rows x 125 columns



```
y=dfe['FUEL CONSUMPTION']
```

y

| | |
|-------|----|
| 0 | 55 |
| 1 | 48 |
| 2 | 85 |
| 3 | 97 |
| 4 | 63 |
| ... | .. |
| 22551 | 70 |
| 22552 | 68 |
| 22553 | 73 |
| 22554 | 78 |
| 22555 | 87 |

Name: FUEL CONSUMPTION, Length: 22556, dtype: int64

x.dtypes

```
YEAR int64
ENGINE SIZE int64
CYLINDERS int64
FUEL int64
HWY (L/100 km) int64
...
VEHICLE CLASS_TWO-SEATER int64
VEHICLE CLASS_Two-seater int64
VEHICLE CLASS_VAN - CARGO int64
VEHICLE CLASS_VAN - PASSENGER int64
VEHICLE CLASS_Van: Passenger int64
Length: 125, dtype: object
```



```
y.dtypes
```

```
dtype('int64')
```

```
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
test=SelectKBest(score_func=chi2)
fi=test.fit(x, y)
fi.scores_

array([4.62201490e+03, 1.15738655e+05, 7.91521885e+03, 8.41761273e+02,
1.96701058e+05, 2.55445570e+05, 7.23715384e+04, 4.66631266e+05,
2.91324966e+02, 1.37664405e+03, 5.44122214e+02, 2.63084816e+02,
3.82369530e+02, 6.99498270e+02, 4.04658354e+02, 3.02128265e+03,
8.69811221e+02, 8.35107407e+03, 2.72707918e+02, 2.68676072e+03,
1.59592159e+04, 2.21538241e+02, 5.12826005e+02, 1.43557778e+03,
5.25454161e+02, 1.95214048e+02, 5.43075357e+02, 7.53222227e+02,
2.32527593e+02, 1.57400714e+03, 1.74574252e+03, 6.05277770e+03,
8.66329831e+02, 5.58495736e+02, 6.11160184e+02, 1.58250981e+02,
2.33227404e+03, 1.94050226e+02, 1.60774817e+03, 9.60718662e+02,
6.21125111e+02, 1.45831096e+03, 1.34523296e+03, 3.53164451e+02,
2.13023297e+02, 7.13352398e+02, 8.84180470e+02, 5.77172506e+02,
4.60180945e+02, 6.31118154e+02, 5.47247991e+02, 9.90635449e+02,
3.77104109e+03, 1.28751934e+03, 6.46061149e+02, 4.61876983e+02,
8.06228215e+03, 2.64193757e+02, 8.54793124e+02, 3.35706534e+02,
1.72886752e+03, 5.30722821e+02, 7.38870523e+02, 2.05126656e+03,
3.87748853e+02, 7.06955129e+02, 7.41136349e+02, 3.53253501e+02,
9.03802478e+02, 5.08956175e+02, 4.37301699e+02, 1.45723097e+02,
1.54638889e+02, 3.31523575e+02, 6.59131407e+02, 6.45260734e+02,
4.99450930e+02, 3.20303844e+03, 5.72308579e+02, 4.48421615e+03,
3.60131415e+02, 2.84043299e+02, 8.32115028e+02, 6.87484991e+03,
8.55276190e+02, 1.04262765e+03, 8.09694320e+02, 4.35947978e+02,
1.09856609e+03, 1.62228633e+03, 1.69265342e+03, 6.62306138e+02,
3.00895520e+02, 2.39286651e+02, 7.01619784e+02, 8.31157593e+02,
1.29669888e+03, 8.27889883e+02, 6.39277147e+02, 1.14182518e+03,
1.08325196e+03, 3.89571341e+02, 4.65029894e+02, 6.50734762e+02,
2.74660421e+03, 4.66278971e+02, 1.00647022e+03, 2.07036823e+02,
4.77062215e+02, 9.65903539e+02, 6.61323995e+02, 1.17994008e+03,
5.24390401e+02, 1.23074262e+03, 7.19979808e+02, 1.19218136e+03,
4.51943264e+02, 7.03835787e+02, 1.10554845e+03, 3.20852436e+02,
1.33219484e+03, 2.19447994e+03, 3.28365818e+03, 4.34332498e+03,
8.63804240e+02])
```

```
col=x.columns
score=pd.DataFrame({'features':col,'score_chi2':fi.scores_})
score
```

| | features | score_chi2 |  |  |
|-----|-------------------------------|---------------|---|---|
| 0 | YEAR | 4622.014898 | | |
| 1 | ENGINE SIZE | 115738.655416 | | |
| 2 | CYLINDERS | 7915.218846 | | |
| 3 | FUEL | 841.761273 | | |
| 4 | HWY (L/100 km) | 196701.058097 | | |
| ... | ... | ... | | |
| 120 | VEHICLE CLASS_TWO-SEATER | 1332.194837 | | |
| 121 | VEHICLE CLASS_Two-seater | 2194.479936 | | |
| 122 | VEHICLE CLASS_VAN - CARGO | 3283.658181 | | |
| 123 | VEHICLE CLASS_VAN - PASSENGER | 4343.324980 | | |
| 124 | VEHICLE CLASS_Van: Passenger | 863.804240 | | |

125 rows × 2 columns

```
score.sort_values(by='score_chi2',ascending=False)
```

| | features | score_chi2 |
|-----|--------------------|---------------|
| 7 | EMISSIONS | 466631.266177 |
| 5 | COMB (L/100 km) | 255445.570099 |
| 4 | HWY (L/100 km) | 196701.058097 |
| 1 | ENGINE SIZE | 115738.655416 |
| 6 | COMB (mpg) | 72371.538425 |
| ... | ... | ... |
| 25 | Company_Cadillac | 195.214048 |
| 37 | Company_Genesis | 194.050226 |
| 35 | Company_GENESIS | 158.250981 |
| 72 | Company_PLYMOUTH | 154.638889 |
| 71 | Company_OLDSMOBILE | 145.723097 |

125 rows × 2 columns

```
#train test split
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
```

```
x_train.shape

(18044, 125)
```

```
x_test.shape

(4512, 125)
```

```
#normalize
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train=sc.fit_transform(x_train)
x_test=sc.fit_transform(x_test)
```

```
import pandas as pd
from sklearn.linear_model import LinearRegression, Lasso
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score, mean_squared_error
```

```
#Linear Regression
linear_model=LinearRegression()
linear_model.fit(x_train, y_train)
linear_y_pred=linear_model.predict(x_test)
linear_mse=mean_squared_error(y_test, linear_y_pred)
linear_r2=r2_score(y_test, linear_y_pred)
#Lasso Regression
lasso_model=Lasso(alpha=0.1)
lasso_model.fit(x_train, y_train)
lasso_y_pred=lasso_model.predict(x_test)
lasso_mse=mean_squared_error(y_test, lasso_y_pred)
lasso_r2=r2_score(y_test, lasso_y_pred)
#Random Forest Regressor
rf_model=RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(x_train, y_train)
rf_y_pred=rf_model.predict(x_test)
rf_mse=mean_squared_error(y_test, rf_y_pred)
rf_r2=r2_score(y_test, rf_y_pred)
#Decision Tree Regressor
dt_model=DecisionTreeRegressor(random_state=42)
dt_model.fit(x_train, y_train)
dt_y_pred=dt_model.predict(x_test)
dt_mse=mean_squared_error(y_test, dt_y_pred)
dt_r2=r2_score(y_test, dt_y_pred)
```

```
#compare the MSE values to see which regressor has the lowest MSE
min_mse=min(linear_mse, lasso_mse, rf_mse, dt_mse)
```

```
best_model=None
if min_mse==linear_mse:
    best_model="Logistic Regression"
elif min_mse==lasso_mse:
    best_model="Lasso Regression"
elif min_mse==rf_mse:
    best_model="Random Forest Regression"
else:
    best_model="Random Forest Regressor"
print("\nThe best performing model is:", best_model, "with mse", linear_mse)
```

The best performing model is: Logistic Regression with mse 1.3479226877777626

```
#print the r2 score results
print("Linear Regression")
print("LR r2 score:", linear_r2)
print("\nLasso Regression:")
print("L r2 score:", lasso_r2)
print("\nRandom Forest Regressor:")
print("RF r2 score:", rf_r2)
print("\nDecision Tree Regressor:")
print("DT r2 score:", dt_r2)
```

Linear Regression
LR r2 score: 0.9988681246786404

Lasso Regression:
L r2 score: 0.9980411787536116

Random Forest Regressor:
RF r2 score: 0.9980340943871755

Decision Tree Regressor:
DT r2 score: 0.9973224217428676

```
#compare the R2 scores to see which regressor has the highest R2 score.
max_r2=max(linear_r2, lasso_r2, rf_r2, dt_r2)
best_model=None
if max_r2==linear_r2:
    best_model="Linear Regression"
elif max_r2==lasso_r2:
    best_model="Lasso Regression"
else:
    best_model="Random Forest Regressor"
print("\nThe best performing model is:", best_model, "with r2 score", rf_r2)
```

The best performing model is: Linear Regression with r2 score 0.9980340943871755

```
#finding the difference between y test and linear y pred
results=pd.DataFrame()
results['Actual']=y_test
results['Predicted']=linear_y_pred
results['Difference']=y_test-linear_y_pred
results.sort_index()
```

| | Actual | Predicted | Difference |
|-------|--------|-----------|------------|
| 3 | 97 | 96.292567 | 0.707433 |
| 17 | 97 | 95.616279 | 1.383721 |
| 31 | 93 | 93.171219 | -0.171219 |
| 34 | 78 | 77.196463 | 0.803537 |
| 35 | 86 | 85.633820 | 0.366180 |
| ... | ... | ... | ... |
| 22526 | 22 | 20.704717 | 1.295283 |
| 22532 | 61 | 59.899517 | 1.100483 |
| 22541 | 58 | 55.957742 | 2.042258 |
| 22551 | 70 | 69.927146 | 0.072854 |
| 22554 | 78 | 77.054206 | 0.945794 |

4512 rows × 3 columns

The best performing model is: Linear Regression

✓

0s

completed at 7:19 PM

×

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.